



Mastra agent series: Book 1

Followed by: *Patterns for Building AI Agents*

Rapid advances in large language models (LLMs) have enabled new kinds of AI applications, known as agents. Written by a veteran of web development, this book walks through:

- The key building blocks of agents: providers, models, prompts, tools, memory, sandboxes
- How to break down complex tasks with agentic workflows
- Agent observability with tracing and evals

Understanding frontier tech is essential for building the future. Sam has done it once with Gatsby, and now again with Mastra — **Paul Klein, CEO of Browserbase**

If you're trying to build agents or assistants into your product, you need to read "Principles" ASAP — **Peter Kazanjy, Author of Founding Sales and CEO of Atrium**



Sam Bhagwat is the founder of Mastra, an open-source JavaScript agent framework, and previously the co-founder of Gatsby, the popular React framework.

\$12.99
ISBN 979-8-9939235-0-5
5 1299 >



9 798993 923505


Request more
free copies



Principles of Building AI Agents

Principles of Building AI Agents

3rd Edition



Sam Bhagwat

Sam Bhagwat
Cofounder & CEO Mastra.ai

PRINCIPLES OF BUILDING AI AGENTS

SAM BHAGWAT

CONTENTS

<i>Foreword</i>	ix
Sam Bhagwat	
<i>Introduction</i>	xiii

PART I

PROMPTING A LARGE LANGUAGE MODEL (LLM)

1. A BRIEF HISTORY OF LLMS	3
2. CHOOSING A PROVIDER AND MODEL	5
Hosted vs open-source	5
Model size: Accuracy vs cost/latency	6
Context window size	6
Reasoning models	7
Providers and models (March 2026)	7
3. WRITING GREAT PROMPTS	8
Use the system prompt	8
Give the LLM more examples	9
A “seed crystal” approach	9
Weird formatting tricks	9
Example: A great prompt	10

PART II

BUILDING AN AGENT

4. AGENTS 101	15
Levels of Autonomy	15
Creating a basic agent	16
5. MODEL ROUTING AND STRUCTURED OUTPUT	18
Structured output	19

6. TOOL CALLING	20
Designing your tools: The most important step	22
Real-world example: Alana's book recommendation agent	22
Better approach:	22
Result:	23
Takeaway:	23
7. AGENT MEMORY	24
Observational memory	25
Memory processors	26
Prompt caching	28
8. DYNAMIC AGENTS	29
What are Dynamic Agents?	29
Example: Creating a Dynamic Agent	30
9. AGENT MIDDLEWARE	31
Guardrails	31
Agent authentication and authorization	33

PART III

TOOLS & MCP

10. MODEL CONTEXT PROTOCOL (MCP): CONNECTING AGENTS AND TOOLS	37
What is MCP?	37
MCP Primitives	38
The MCP Ecosystem	38
When to use MCP	39
Building an MCP Server and Client	39
11. POPULAR THIRD-PARTY TOOLS FOR AGENTS	42
Web scraping & computer use	42
Third-party integrations	43

PART IV
GRAPH-BASED WORKFLOWS

12. WORKFLOWS 101	47
13. BRANCHING, CHAINING, MERGING, CONDITIONS	48
Branching	48
Chaining	49
Merging	50
Conditions	50
Best practices	51
14. SUSPEND AND RESUME	52
15. STREAMING UPDATES	56
Streaming step completion	57
Streaming within steps	57
Streaming tool calls	58
Speed matters	58

PART V
**RETRIEVAL-AUGMENTED
GENERATION (RAG)**

16. RAG 101	61
17. CHOOSING A VECTOR DATABASE	63
18. SETTING UP YOUR RAG PIPELINE	65
Chunking	65
Embedding	66
Upsert	66
Indexing	66
Querying	66
Reranking	67
Code Example	67
19. ALTERNATIVES TO RAG	70
Give your agent search tools	70
Let your agent run code	71
Feed the model full context	71
Extract entities and relationships	71
Conclusion	71

PART VI
MULTI-AGENT SYSTEMS

20. MULTI-AGENT 101	75
21. AGENT SUPERVISOR & SUBAGENTS	77
22. CONTROL FLOW	78
23. WORKFLOWS AS TOOLS	79
24. PARALLELIZED TOOL CALLS	80
25. COMBINING THE PATTERNS	81
A note on A2A	81

PART VII
OBSERVABILITY & EVALS

26. TRACING	85
Why accuracy and token cost matter	85
Observability	86
Visualizing traces	86
Setting this up on a project	87
27. EVALS	89
Seeing evals over time	90
Building your eval dataset	91
LLM-as-judge for textual evals	92
Classification or labeling evals	93
Tool calling evals	93
Multi-turn evals	93
Task completion	93
Prompt engineering evals	94
A/B testing	94
Human data review	94
Observability & evals are even more important than you think	95

PART VIII
DEVELOPMENT & DEPLOYMENT

28. LOCAL DEVELOPMENT	99
Building an agentic web frontend	99
Building an agent backend	100

29. DEPLOYMENT	102
The core agent loop	102
Agentic workflows	102
Using a managed platform	103

PART IX

CODING AGENTS

30. AGENTS BUILDING AGENTS	107
On web-based platforms	107
31. SANDBOXES & FILESYSTEMS	109
Ephemeral vs. stateful sandboxes	109
Filesystems	110
Managing latency	110

PART X

EVERYTHING ELSE

32. MULTIMODAL	113
Image Generation	114
Use Cases	115
Voice	115
Video	117
33. WHAT'S NEXT	118

FOREWORD

SAM BHAGWAT

3rd edition

Three editions of one book in one year? One thing we've learned at Mastra: time moves faster in the AI world.

This new edition includes new material on coding agents. There are also text, image, and code sample updates throughout the text to reflect what the AI agent space looks like in March 2026.

In November we also released a companion book, *Patterns for Building AI Agents*. In the last six months we've seen a lot of folks putting their agents into production with Mastra, and so *Patterns* is a collection of strategies and lessons learned from advanced practitioners, with deep-dives on context engineering, evals, and security.

I also realized in the rush to get the first couple editions out the door we forgot the thank yous and dedications. This book would not have been possible without my cofounders, Shane and Abhi, and my partner, Jenn.

It's dedicated to my children, Oliver (8), and Hazel (6),

who are growing up in a world where their conversations are not just with people, but with models and agents.

Sam Bhagwat
San Francisco, CA
March 2026

2nd edition

Two months is a short time to write a new edition of a book, but life moves fast in AI.

This edition has new content on MCP, image gen, voice, A2A, web browsing and computer use, workflow streaming, code generation, agentic RAG, and deployment.

AI engineering continues to get hotter and hotter. Mastra's weekly downloads have doubled each of the last two months. At a typical SF AI evening meetup, I give away a hundred copies of this book.

Then two days ago, a popular developer newsletter tweeted about this book and 3,500 people (!) downloaded a digital copy (available for free at mastra.ai/book if you are reading a paper copy).

So yes, 2025 is truly the year of agents. Thanks for reading, and happy building!

Sam Bhagwat
San Francisco, CA
May 2025

1st edition

For the last three months, I've been holed up in an apartment in San Francisco's Dogpatch district with my cofounders, Shane Thomas and Abhi Aiyr.

We're building an open-source JavaScript framework called Mastra to help people build their own AI agents and assistants.

We've come to the right spot.

We're in the Winter 2025 batch of YCombinator, the most popular startup incubator in the world (colloquially, YC W25).

Over half of the batch is building some sort of "vertical agent" — AI application generating CAD diagrams for aerospace engineers, Excel financials for private equity, a customer support agent for iOS apps.

These three months have come at some personal sacrifice.

Shane has traveled from South Dakota with his girlfriend Elizabeth, their three-year-old daughter and newborn son. I usually have 50-50 custody of my seven-year-old son and five-year-old daughter, but for these three months I'm down to every-other-weekend. Abhi's up from LA, where he bleeds Lakers purple and gold.

Our backstory is that Shane, Abhi and I met while building a popular open-source JavaScript website framework called Gatsby. I was the co-founder, and Shane and Abhi were key engineers.

While OpenAI and Anthropic's models are widely available, the secrets of building effective AI applications are hidden in niche Twitter/X accounts, in-person SF meetups, and founder groupchats.

But AI engineering is just a new domain, like data engi-

neering a few years ago, or DevOps before that. It's not impossibly complex. An engineer with a framework like Mastra should be able to get up to speed in a day or two. With the right tools, it's easy to build an agent as it is to build a website.

This book is intentionally a short read, even with the code examples and diagrams we've included. It should fit in your back pocket, or slide into your purse. You should be able to use the code examples and get something simple working in a day or two.

Sam Bhagwat

San Francisco, CA

March 2025

INTRODUCTION

We've structured this book into a few different sections.

Prompting a Large Language Model (LLM) provides some background on what LLMs are, how to choose one, and how to talk to them.

Building an Agent introduces a key building block of AI development. Agents are a layer on top of LLMs: they can execute code, store and access memory, and communicate with other agents. Chatbots are typically powered by agents.

Graph-based Workflows have emerged as a useful technique for building with LLMs when agents don't deliver predictable enough output.

Retrieval-Augmented Generation (RAG), covers a common pattern of LLM-driven search. RAG helps you search through large corpuses of (typically proprietary) information in order to send the relevant bits to any particular LLM call.

Multi-Agent Systems cover the coordination aspects of bringing agents into production. The problems often involve a significant amount of organizational design!

Testing with Evals is important in checking whether your application is delivering users sufficient quality.

Local Development and Deployment are the two places where your code needs to work. You need to be able to iterate quickly on your machine, then get code live on the Internet.

Agents Building Agents is about the ways building agents changes in an era where we're increasingly writing code via tools like Claude Code rather than by hand.



Note that we don't talk about traditional machine learning (ML) topics like reinforcement learning, training models, and fine-tuning.

Today most AI applications only need to *use* LLMs, rather than *build them*.

PART I

PROMPTING A LARGE LANGUAGE MODEL (LLM)

A BRIEF HISTORY OF LLMs

For over 40 years, AI was a perennial on-the-horizon technology.

There have been some notable advances along the way, like chess engines, speech recognition, self-driving cars.

But the true turning point for progress on “generative AI” was 2017, when eight researchers from Google wrote a paper called “Attention is All You Need”.

It described an architecture for generating text where a “large language model” (known as an “LLM” or “model”) could be given a set of “tokens” (words and punctuation) and focused on predicting the next “token”.

Then a great leap forward happened in November 2022. A chat interface called ChatGPT, produced by a well-funded startup called OpenAI, went viral literally overnight by making generative AI accessible to everyday users.

Today, LLMs are proliferating at an unbelievable rate. Major providers, which provide both consumer chat interfaces and developer APIs, include:

- **OpenAI.** Founded in 2015 by eight people including AI researcher Ilya Sutskever, software engineer Greg Brockman, Sam Altman (the head of YC), and Elon Musk.
- **Anthropic (Claude).** Founded in 2020 by Dario Amodei and a group of former OpenAI researchers. Produces models popular for code writing, as well as API-driven tasks.
- **Google (Gemini).** The core LLM is being produced by the DeepMind team acquired by Google in 2014.
- **Meta (Llama).** The Facebook parent company produces the Llama class of open-source models. Once considered the leading US open-source AI group.
- **Others** include Mistral (an open-source French company), DeepSeek (an open-source Chinese company) and Qwen (open-source, from Alibaba Cloud in China).

CHOOSING A PROVIDER AND MODEL

One of the first choices you'll need to make building an AI application is which model to build on.

Here are some considerations:

Hosted vs open-source

The first piece of advice we usually give people when building AI applications is to start with a hosted provider.

For text modality use cases — including text, code, structured output, tool calls (but not images, voice, or video), the state-of-the-art hosted model providers in March 2026 are OpenAI, Anthropic, or Google Gemini. We recommend you start there.

Perhaps you eventually move some of your workload to SoTA open-source providers like Qwen, Kimi, Deepseek, Minimax. But *make it work, make it right, make it fast/cheap* — in that order. Don't make yourself debug infra issues right off the bat.

To give yourself flexibility later, use a model routing library (more on that later).

Model size: Accuracy vs cost/latency

Every provider has larger, medium, and small models (Anthropic's Claude 4.6 Opus vs Sonnet vs Haiku, OpenAI's GPT-5, GPT-5 mini, GPT-5 nano).

The larger a model is, the more mathematical operations (matrix multiplication) it has to perform to generate your answer. That makes it higher quality, but also more expensive and slower.

Again, *make it work, make it right, make it fast/cheap* — in that order.

Start with the larger models, when you need to worry about cost saving during scaling, then figure out how to move workloads to cheaper models.

Context window size

One variable you may want to think about is the “context window” of your model. How many tokens can it take?

As of March 2026, the longest context window is offered by Google. Gemini 3.1 supports a one million token context window (roughly 2,000 pages of text), with Anthropic offering the same for Sonnet 4.6 (but only in the API, in beta)

This allows some interesting applications; we know a developer who built a support assistant for an iOS app and simply dumped the entire 800,000-token codebase in its context window.

Reasoning models

Another type of model is what’s called a “reasoning model.” Starting in late 2024, models began to do a lot of logic internally before returning a response.

It might take seconds, or minutes, to give a response, and it will return a response all at once (while streaming some “thinking steps” along the way).

Today most state-of-the-art models like GPT 5.4 are reasoning models — though most let you can turn down effort to low when using them in an API to reduce latency. For lower-complexity tasks, you could also use pre-reasoning models like GPT-4o.

Providers and models (March 2026)

Provider	Form Factor(s)	Default Model	Cheap / Fast Model	Reasoning Model(s)
OpenAI	Hosted	GPT-5.4	GPT-5 mini, GPT-5 nano	Reasoning built into GPT-5.4 with effort controls; GPT-5.2 Thinking / Pro
Anthropic	Hosted	Claude Sonnet 4.6	Claude Haiku 4.5	Claude Opus 4.6
Google Gemini	Hosted	Gemini 3 Flash	Gemini 3.1 Flash-Lite	Gemini 3.1 Pro
DeepSeek	OSS	DeepSeek V3.2	Efficient by default (MoE architecture)	DeepSeek R1 series
Alibaba (Qwen)	OSS	Qwen 3.5	Qwen3-30B-A3B	QwQ-32B

WRITING GREAT PROMPTS

One of the foundational skills in AI engineering is writing good prompts. LLMs will follow instructions, if you know how to specify them well. Here's a few tips and techniques that will help:

Use the system prompt

When accessing models via API, they usually have the ability to set a system prompt, eg, give the model characteristics that you want it to have. This will be in addition to the specific “user prompt” that gets passed in.

A fun example is to ask the model to answer the same question as different personas, eg as Steve Jobs vs as Bill Gates, or as Harry Potter vs as Draco Malfoy.

Beyond tone, you can also use the system prompt to improve accuracy.

Giving the model a clear role ("You are an expert tax attorney"), specific constraints ("Always cite your sources"), or explicit instructions about what to avoid ("Do not make up information you are not certain

about") can meaningfully reduce errors and hallucinations.

Give the LLM more examples

There are three basic techniques to prompting.

- **Zero-shot:** The “YOLO” approach. Ask the question and hope for the best.
- **Single-shot:** Ask a question, then provide one example (w/ input + output) to guide the model
- **Few-shot:** Give **multiple** examples for more precise control over the output.

More examples = more guidance, but this also takes more time.

One consideration is prompt caching, a feature introduced across most providers in 2025. Putting examples in a cached system prompt rather than the user turn is now standard practice for both cost and latency reasons.

A “seed crystal” approach

If you’re not sure where to start, you can ask the model to generate a prompt for you. E.g. “Generate a prompt for requesting a picture of a dog playing with a whale.”

We actually built a prompt CMS into Mastra’s local development environment for this reason.

Weird formatting tricks

AI models can be sensitive to formatting—use it to your advantage. While ALL CAPS for emphasis is a relic of

models like GPT-3.5, XML-like structure or markdown can help models follow instructions more precisely.

Different models respond differently. Claude tends to follow formatting when given structural scaffolding like XML-style tags. GPT responds better to markdown-style syntax and delimiter cues.

Example: A great prompt

If you think your prompts are detailed, go through and read some production prompts. Here's about one-third of a live production code-generation prompt (used in a tool called bolt.new).

You are Bolt, an expert AI assistant and exceptional senior software developer with vast knowledge across multiple programming languages, frameworks, and best practices.

<system_constraints>

You are operating in an environment called WebContainer, an in-browser Node.js runtime that emulates a Linux system to some degree. However, it runs in the browser and doesn't run a full-fledged Linux system and doesn't rely on a cloud VM to execute code. All code is executed in the browser. It does come with a shell that emulates zsh. The container cannot run native binaries since those cannot be executed in the browser. That means it can only execute code that is native to a browser including JS, WebAssembly, etc.

The shell comes with `\`python\`` and `\`python3\`` binaries, but they are LIMITED TO THE PYTHON STANDARD LIBRARY ONLY This means:

- There is NO `\`pip\`` support! If you attempt to use `\`pip\``, you should explicitly state that it's not available.
- CRITICAL: Third-party libraries cannot be installed or imported.
- Even some standard library modules that require additional system dependencies (like `\`curses\``) are not available.
- Only modules from the core Python standard library can be used.

Additionally, there is no `\`g++\`` or any C/C++ compiler available. WebContainer CANNOT run native binaries or compile C/C++ code!

Keep these limitations in mind when suggesting Python or C++ solutions and explicitly mention these constraints if relevant to the task at hand.

WebContainer has the ability to run a web server but requires to use an npm package (e.g., Vite, server, serve, http-server) or use the Node.js APIs to implement a web server.

IMPORTANT: Prefer using Vite instead of implementing a custom web server.

IMPORTANT: Git is NOT available.

IMPORTANT: Prefer writing Node.js scripts instead of shell scripts. The environment doesn't fully support shell scripts, so use Node.js for scripting tasks whenever possible!

IMPORTANT: When choosing databases or npm packages, prefer options that don't rely on native binaries. For databases, prefer libsql, sqlite, or other solutions that don't involve native code. WebContainer CANNOT execute arbitrary native binaries.

Available shell commands: cat, chmod, cp, echo, hostname, kill, ln, ls, mkdir, mv, ps, pwd, rm, rmdir, xxd, alias, cd, clear, curl, env, false, getconf, head, sort, tail, touch, true, uptime, which, code, jq, loadenv, node, python3, wasm, xdg-open, command, exit, export, source

</system_constraints>

PART II

BUILDING AN AGENT

AGENTS 101

You can use direct model calls for one-shot transformations: “*Given a video transcript, write a draft description.*”

But for ongoing, complex interactions with models, you typically need to build an agent on top.

Think of agents as AI employees rather than contractors: they maintain context, have specific roles, and can use tools to accomplish tasks.

Levels of Autonomy

In September 2025 Simon Willison distilled a fundamental definition of agents that we think is pretty solid: *An agent calls tools in a loop to achieve a goal.*

Agency is a spectrum. Like self-driving cars, there are different levels of agent autonomy.

- At a low level, agents make binary choices in a decision tree

- At a medium level, agents have memory, call tools, and retry failed tasks
- At a high level, agents do planning, divide tasks into subtasks, manage their task queue, manage multiple parallel sub-agents, and self-correct across long task horizons.

Creating a basic agent

In Mastra, agents have persistent memory, consistent model configurations, and can access a suite of tools and workflows.

The **agent** class is the foundation for creating AI agents in Mastra. It provides methods for generating responses, streaming interactions, and handling voice capabilities.

Here's how to create a basic agent:



```
import { Memory } from '@mastra/memory'
import { Agent } from '@mastra/core/agent'

export const agent = new Agent({
  name: 'my-agent',
  instructions: 'You are a helpful assistant',
  model: 'openai/gpt-5.1',
})
```

Running this code creates a configured AI agent that can:

- Accept user inputs/conversations
- Respond according to the "helpful assistant" instructions
- Use GPT-5.1 as the underlying language model
- Be integrated into larger applications or workflows

It's worth noting that unlike a single model call, agents accumulate tokens quickly. Tool results, memory, and conversation history all add up.

This make context management and cost key things you'll need to figure out — not necessarily when you're initially building, but certainly as you begin rolling out your agent to users.

MODEL ROUTING AND STRUCTURED OUTPUT

It's useful to be able to quickly test and experiment with different models. Various libraries provide an abstraction known as *model routing*.

Let's pull up that example again from last chapter, which uses Mastra's built-in model routing.

```
import { Memory } from '@mastra/memory'
import { Agent } from '@mastra/core/agent'

export const agent = new Agent({
  name: 'my-agent',
  instructions: 'You are a helpful assistant',
  model: 'openai/gpt-5.1',
})

const result = await agent.generate("What is the weather today?");
```

The nice thing about structuring the code this way, whether you're using Mastra or something else, is that it's a one-line fix to swap in a different model and provider — instead of needing to rip out a provider SDK.

Structured output

When you use models as part of an application, you often want them to return data in JSON format instead of unstructured text. Most models support “structured output” to enable this.

Here’s an example of requesting a structured response by providing a schema:



```
import { z } from 'zod'

const response = await agent.generate('Help me plan my day.', {
  structuredOutput: {
    schema: z.array(
      z.object({
        name: z.string(),
        activities: z.array(z.string()),
      })
    ),
  },
})
```

Models are very powerful for processing unstructured or semi-structured text. Consider passing in the text of a resume and extracting a list of jobs, employers, and date ranges, or passing in a medical record and extracting a list of symptoms.

Structured output and model routing actually compose nicely together. You can usually route to a cheaper/faster model specifically for structured extraction tasks, since they tend to require less reasoning than open-ended generation.

TOOL CALLING

Tools are functions that agents can call to perform specific tasks — whether that's fetching weather data, querying a database, or processing calculations.

The key to effective tool use is clear communication with the model about what each tool does and when to use it.

Here's an example of creating and using a tool:

```
import { createTool } from "@mastra/core/tools";
import { z } from "zod";

const getWeatherInfo = async (city: string) => {
  // Replace with an actual API call to a weather
  service
  console.log(`Fetching weather for ${city}...`);
  // Example data structure
  return { temperature: 20, conditions: "Sunny" };
};

export const weatherTool = createTool({
  id: "Get Weather Information",
  description: `Fetches the current weather
  information for a given city`,
  inputSchema: z.object({
    city: z.string().describe("City name"),
  }),
  outputSchema: z.object({
    temperature: z.number(),
    conditions: z.string(),
  }),
  execute: async ({ context: { city } }) => {
    console.log("Using tool to fetch weather
    information for", city);
    return await getWeatherInfo(city);
  },
});
```

Best practices:

- Provide detailed descriptions in the tool definition and system prompt
- Use specific input/output schemas
- Use semantic naming that matches the tool's function (eg **multiplyNumbers** instead of **doStuff**)

Remember: The more clearly you communicate a tool's purpose and usage to the model, the more likely it is to use it correctly. You should describe both what it does and when to call it.

Designing your tools: The most important step

When you're creating an AI application, the most important thing you should do is think carefully about your tool design:

- What the tools will you need?
- What will each of them do?

It is genuinely important that you *write this list out clearly before you start coding*, even if it's just jotted down on a napkin left over from lunch at your desk.

Real-world example: Alana's book recommendation agent

Alana Goyal, a Mastra investor, wanted to build an agent that could analyze, and give intelligent recommendations about, her library of investor books.

First attempt:

She tried dropping all the books into the agent's knowledge window. This didn't work well — the agent couldn't reason about the data in a structured way.

Better approach:

She broke the problem down into a set of specific tools, each handling a different aspect of the data:

- A tool for accessing the corpus of investors

- A tool for book recommendations
- A tool for books tagged by genre

Then, she added more tools for common operations:

- Get all books by genre
- Get book recommendations by investor
- Sort people writing recommendations by type (founders, investors, etc.)

If a human analyst were doing this project, they'd follow a specific set of operations or queries.

The trick is to take those operations and write them as tools or queries that your agent can use.

Result:

With these tools in place, the agent could now intelligently analyze the corpus of books, answer nuanced questions, and provide useful recommendations — just like a skilled human analyst.

Takeaway:

Think like an analyst. Break your problem into clear, reusable operations. Write each as a tool.

If you do this, your agent will be much more capable, reliable, and useful.

AGENT MEMORY

Memory is crucial for creating agents that maintain meaningful, contextual conversations over time. While LLMs can process effectively individual messages, they need help managing longer-term context and historical interactions.

There are three established types of agent memory:

- “working memory” (relevant long-term characteristics of users, eg how ChatGPT store what it knows about you)
- “semantic recall” (embedding text as vectors and using semantic search, or RAG, to find similar results to the query string).
- “observational memory” — a more agentic approach that compresses sessions into structured observations."

Observational memory

If you step outside on a busy street, your brain processes millions of pixels, but distills down to one or two observations. That blue SUV just ran a red light. Your neighbor's pit bull is off their leash.

In the context of a coding agent, observational memory might compress a user session down to something like this:

Date: 2026-01-15

- 🟡 12:10 User is building a Next.js app with Supabase auth, due in 1 week (meaning January 22nd 2026)
- 🟡 12:10 App uses server components with client-side hydration
- 🟡 12:12 User asked about middleware configuration for protected routes
- 🟡 12:15 User stated the app name is "Acme Dashboard"

The core message format uses **formatted text, not structured objects**. (text is the universal interface and it's optimized for LLMs), a **three-date model** (observation date, referenced date, and relative date) and **emoji based log levels**.

In observational memory, the context window is broken into two blocks. The first block is the list of observations (like above). The second block is raw messages that haven't yet been compressed. Typically 30k tokens are reserved for the first block and 40k for the second block.

When new messages come in, they are appended to the

end of the second block. When raw messages overflow, a separate "observer agent" compresses messages into new observations; when observations overflow a separate "reflector agent" garbage collects.



```
import { Memory } from '@mastra/memory'
import { Agent } from '@mastra/core/agent'

export const agent = new Agent({
  name: 'my-agent',
  instructions: 'You are a helpful assistant',
  model: 'openai/gpt-5-codex',
  memory: new Memory({
    options: {
      observationalMemory: true,
    },
  }),
})
```

Memory processors

Sometimes increasing your context window is not the right solution. It may seem counterintuitive, but sometimes you want to deliberately prune your context window (or just have more control over it).

Memory processors allow you to modify the list of messages retrieved from memory *before* they are added to the agent's context window and sent to the LLM. This is useful for managing context size, filtering content, and optimizing performance (not mention saving tokens).

Mastra provides some built-in memory processors:

. . .

TokenLimiter

This processor is used to prevent errors caused by exceeding the LLM's context window limit. It counts the tokens in the retrieved memory messages and removes the oldest messages until the total count is below the specified limit.


```
import { Memory } from "@mastra/memory";
import { TokenLimiter } from
"@mastra/memory/processors";
import { Agent } from "@mastra/core/agent";
import { openai } from "@ai-sdk/openai";

const agent = new Agent({
  model: openai("gpt-4o"),
  memory: new Memory({
    processors: [
      // Ensure the total tokens from memory don't
      // exceed ~127k
      new TokenLimiter(127000),
    ],
  }),
});
```

ToolCallFilter

This processor removes tool calls from the memory messages sent to the LLM. It saves tokens by excluding potentially verbose tool interactions from the context, which is useful if the details aren't needed for future interactions. It's also useful if you always want your agent to call

a specific tool again and not rely on previous tool results in memory.



```
import { Memory } from "@mastra/memory";
import { ToolCallFilter, TokenLimiter } from
"@mastra/memory/processors";

const memoryFilteringTools = new Memory({
  processors: [
    // Example 1: Remove all tool calls/results
    new ToolCallFilter(),

    // Example 2: Remove only noisy image generation
    tool calls/results
    new ToolCallFilter({ exclude:
["generateImageTool"] }),

    // Always place TokenLimiter last
    new TokenLimiter(127000),
  ],
});
```

You can also create more custom memory logic by extending the base **MemoryProcessor** class.

Prompt caching

Most major APIs now let you cache a static prefix, like a long system prompt, a knowledge base, or a set of examples.

This reduces token cost significantly while keeping frequently used context available, and plays nicely with observational memory.

DYNAMIC AGENTS

The simplest way to configure an agent is to pass a string for its system prompt, a string for the LLM provider and model name, and an object/dictionary for a list of tools that they are provided.

But what if you want to change these things at runtime?

What are Dynamic Agents?

A dynamic agent is an agent whose properties—like instructions, model, and available tools—can be determined at runtime, not just when the agent is created.

This means your agent can change how it acts based on user input, environment, or any other runtime context you provide.

Choosing between dynamic and static agents is ultimately a tradeoff between predictability and power.

Example: Creating a Dynamic Agent

Here's an example of a dynamic support agent that adjusts its behavior based on the user's subscription tier and language preferences:

```
const supportAgent = new Agent({
  name: "Dynamic Support Agent",

  instructions: async ({ runtimeContext }) => {
    const userTier = runtimeContext.get("user-tier");
    const language = runtimeContext.get("language");

    return `You are a customer support agent for our SaaS platform.
    The current user is on the ${userTier} tier and prefers ${language}
    language.

    For ${userTier} tier users:
    ${userTier === "free" ? "- Provide basic support and documentation
    links" : ""}
    ${userTier === "pro" ? "- Offer detailed technical support and best
    practices" : ""}
    ${userTier === "enterprise" ? "- Provide priority support with
    custom solutions" : ""}

    Always respond in ${language} language.`;
  },

  model: ({ runtimeContext }) => {
    const userTier = runtimeContext.get("user-tier");
    return userTier === "enterprise"
      ? openai("gpt-4")
      : openai("gpt-3.5-turbo");
  },

  tools: ({ runtimeContext }) => {
    const userTier = runtimeContext.get("user-tier");
    const baseTools = [knowledgeBase, ticketSystem];

    if (userTier === "pro" || userTier === "enterprise") {
      baseTools.push(advancedAnalytics);
    }

    if (userTier === "enterprise") {
      baseTools.push(customIntegration);
    }

    return baseTools;
  },
});
```


AGENT MIDDLEWARE

Once we see that it's useful to specify the system prompt, model, and tool options at runtime, we start to think about the other things we might want to do at runtime as well.

Guardrails

Guardrails are a general term for sanitizing the input coming into your agent and/or the output it's returning before the user can see it. Input sanitization tries broadly to guard against prompt injection attacks.


These include model “jailbreaking” (“IGNORE PREVIOUS INSTRUCTIONS AND...”), requests for PII, and off-topic chats that could run up your bills.



Models have gotten better at resisting direct jailbreaks, but prompt injection has gotten more sophisticated as agents have gained more autonomy.

An agent that browses the web or reads uploaded documents can encounter malicious instructions embedded in that content.

For example, a February 2026 prompt injection attack against the open-source coding IDE Cline began with an attacker opening a Github issue in a Cline repo using a title containing malicious instructions.



Chris Bakke prompt injection attack, December 2023

Agent authentication and authorization

One thing to think about when building agents is that because they are more powerful than pre-LLM data access patterns, you may need to spend more time ensuring they are *permissioned accurately*.

Security through obscurity becomes less of a viable option when users can ask an agent to retrieve knowledge hidden in nooks and crannies!

There are two layers of permissions to consider for agents:

- The resources the agent itself can have access to
- Which users are allowed to access the agent

As with guardrails, middleware is the typical place to put any agent authorization because it's in the perimeter *around* the agent rather than within the agent's inner loop.

For high-stakes or irreversible actions, consider adding human-in-the-loop checkpoints — pausing execution to get explicit user confirmation before the agent proceeds

PART III

TOOLS & MCP

MODEL CONTEXT PROTOCOL (MCP): CONNECTING AGENTS AND TOOLS

LLMs, like humans, become much more powerful when given tools. MCP works like a USB-C port for AI applications, providing a standard way for models and tools to interact.

What is MCP?

In November 2024, a small team at Anthropic proposed MCP as a solution to a real and growing problem: every AI provider and tool author had their own way of defining and calling tools.

It's an open protocol for connecting AI agents to tools, models, and each other. Think of it as a universal adapter: if your tool or agent “speaks” MCP, it can plug into any other MCP-compatible system—no matter who built it or what language it's written in.

But, as any experienced engineer knows, the power of any protocol is in the network of people following it.

While initially well-received, it took until March for

MCP to hit critical mass after gaining popularity among prominent, vocal supporters like Shopify's CEO Tobi Lutke.

Then, in April 2025, OpenAI and Google Gemini announced they would support MCP, making it the functional default for connecting AI agents. At this point (March 2026) MCP is the established standard for connecting AI agents to third-party tools.

MCP Primitives

MCP has two basic primitives: *servers* and *clients*.

Servers wrap sets of MCP tools. They (and their underlying tools) can be written in any language and communicate with clients over stdio (locally) and Streamable HTTP (for remote servers).

Clients such as models or agents can query servers to get the set of tools provided, then request that the server execute a tool and return a response.

As such, MCP is as a standard for remote code execution, like OpenAPI or RPC.

The MCP Ecosystem

As MCP was gaining traction, a bunch of folks joined the fray.

- *Vendors* like Stripe began shipping MCP servers for their API functionality.
- *Independent developers* started making MCP servers for functionality they needed, like browser use or error tracking, and publishing them on Github

- *Frameworks* like Mastra started shipping MCP server and client abstractions so that individual developers didn't have to reimplement specs themselves.

When to use MCP

Agents, like SaaS, often need a number of basic integrations with third-party services (calendar, chat, email, web). If your project roadmap has a lot of this type of features, it's worth looking at building an MCP *client* that could access third-party tools and services.

Conversely, if you're building a tool that you want *other* agents to use, you should consider shipping an MCP *server*.

When consuming third-party MCP servers, treat them with the same trust level as any third-party API — vet them carefully, especially for agents with access to sensitive data or the ability to take real-world actions.

Building an MCP Server and Client

If you want to create MCP servers and give an agent access to them, here's how you can do that in Typescript with Mastra:

```

// --- weatherTool.ts ---
import { defineTool } from "@mastra/core/tool";

export const weatherTool = defineTool({
  name: "weatherTool",
  description: "Get the current weather for a city.",
  parameters: {
    city: { type: "string", description: "City name" },
  },
  async execute({ city }) {
    // Dummy implementation
    return `The weather in ${city} is sunny!`;
  },
});

// --- weather-server.ts ---
import { MCPServer } from "@mastra/mcp";
import { weatherTool } from "../weatherTool";

const server = new MCPServer({
  name: "Weather Server",
  version: "1.0.0",
  tools: { weatherTool },
});

await server.startStdio();

// --- agent.ts ---
import { MCPClient } from "@mastra/mcp";
import { Agent } from "@mastra/core/agent";
import { openai } from "@ai-sdk/openai";

const mcp = new MCPClient({
  servers: {
    weather: {
      command: "npx",
      args: ["tsx", "weather-server.ts"],
    },
  },
  timeout: 30000,
});

const agent = new Agent({
  name: "Weather Agent",
  instructions: "You can answer weather questions using the weather tool.",
  model: openai("gpt-4"),
  tools: await mcp.getTools(),
});

```

Conversely, if you want to create a client with access to other MCP servers, here's how you would do that:



```
import { MCPClient, MCPServer } from "@mastra/mcp";

// Step 1: Create an MCP client that connects to other MCP servers
const mcp = new MCPClient({
  servers: {
    weather: {
      // Connect to a remote MCP server via HTTP/SSE
      url: new URL("http://localhost:1234/sse"),
    },
    stocks: {
      // Or connect to a local MCP server via stdio
      command: "npx",
      args: ["tsx", "stock-server.ts"],
    },
  },
  timeout: 30000,
});

// Step 2: Expose all tools from the connected MCP servers via a new
// MCPServer
const server = new MCPServer({
  name: "Proxy MCP Server",
  version: "1.0.0",
  tools: await mcp.getTools(), // Aggregate tools from all connected servers
});

// Step 3: Start the proxy MCP server (stdio)
await server.startStdio();
```

POPULAR THIRD-PARTY TOOLS FOR AGENTS

Agents are only as powerful as the tools you give them. As a result, an ecosystem has sprung up around popular types of tools.

Web scraping & computer use

One of the core tool use cases for agents is browser search.

This includes web scraping, automating browser tasks, and extracting information. You can use built-in tools, connect to MCP servers, or integrate with higher-level automation platforms.

There are a few different tools you could take to add search to your agents:

- **Cloud-based web search APIs.** There are a few web search APIs that have become popular for LLMs to use, including Exa, Browserbase, and Tavily.
- **Agentic web search.** Tools like Stagehand (in JavaScript) and Browser Use (in Python, with

MCP servers for JS users) have plain English language APIs that you can use to describe web scraping tasks.

- **Low-level open-source search tools.** Microsoft's Playwright project remains widely used as the underlying engine that many agentic browser tools are built on.

When you provide browser tools to agents, you often encounter similar challenges to traditional browser automation.

- **Anti-bot detection.** From browser fingerprinting to WAFs to captchas, many websites protect against automated traffic.
- **Fragile setups.** Browser use setups sometimes break if target websites change their layout or modify some CSS.

These challenges are solvable — just budget a bit of time for some munging and glue work!

Third-party integrations

Agents need to connect with the systems where user data lives — and they need the ability to both read and write from those systems.

Most agents (like most SaaS apps) need access to a core set of *general* integrations like email, calendar, or documents.

Your personal assistant agent will need access to Gmail, Google Calendar, or Microsoft Outlook.

Your sales agent will need to integrate with Salesforce or

Hubspot. Your HR agent will need to integrate with Rippling or Workday.

It's possible the services you need already ship in an MCP server before you reach for an iPaaS. The ecosystem has grown significantly and many popular tools now offer native MCP support.

Most people building agents want to avoid spending months building bog-standard integrations and choose an “agentic iPaaS” (integration-platform-as-a-service). We've seen folks be happy with Arcade and Composio, which are general services, but there are specialized services in various domains as well.

PART IV

GRAPH-BASED WORKFLOWS

WORKFLOWS 101

We've seen how individual agents can work: Agents have the flexibility to call any tool (function) at every step.

Sometimes, this is *too* much flexibility.

Graph-based workflows have emerged as a more deterministic technique for building with models.

Workflow graphs outline the sequence of actions, branching logic, and conditions that an agent might follow to achieve a specific goal — the flow of execution and decision-making.

In practice, workflows and agents aren't mutually exclusive. You'll often use both, with agents handling open-ended reasoning steps and workflows providing the deterministic scaffolding around them.

Why use a specific workflow abstraction rather than simply write this in code? Workflow primitives are helpful for defining branching logic, parallel execution, creating checkpoints, and adding tracing. Along with a graph for visualization.

Let's dive in.

BRANCHING, CHAINING, MERGING, CONDITIONS

So, what's the best way to build workflow graphs? Several frameworks support workflow graph abstractions. We'll use Mastra here, but the primitives are similar across tools.

Let's walk through the basic operations, and then we can get to best practices.


Branching

One use case for branching is to trigger multiple LLM calls on the same input.

Let's say you have a long medical record, and need to check for the presence of 12 different symptoms (drowsiness, nausea, etc).

You could have one LLM call checks for 12 symptoms. Or, it might be more accurate to have 12 parallel LLM calls, each checking for one symptom.


In Mastra, you create branches with the `.step()` command. Here's a simple example:



Chaining


This is the simplest operation. Sometimes, you'll want to fetch data from a remote source before you feed it into an LLM, or feed the results of one LLM call into another.

In Mastra, you chain with the `.then( )` command. Each step waits for the previous step, and has access to previous step results:



Merging

After branching paths diverge to handle different aspects of a task, they often need to converge again to combine their results:



Conditions

Sometimes your workflow needs to make decisions based on intermediate results.

In workflow graphs, because multiple paths can execute in parallel, in Mastra we define the conditional path execution on the child step rather than the parent step.

In this example, a `processData` step is executing, conditional on the `fetchData` step succeeding.



```
myWorkflow.step(  
  new Step({  
    id: "processData",  
    execute: async ({ context }) => {  
      // Action logic  
    },  
  }),  
  {  
    when: {  
      "fetchData.status": "success",  
    },  
  },  
);
```

Best practices

- Compose steps so that the input/output at each step is meaningful, since you'll be able to see it in your tracing. (More on this in the *Tracing* section a bit further ahead).
- Decompose steps so that the model only has to do one thing at one time. This usually means no more than one model call in any step.

Many different special cases of workflow graphs, like loops, retries, etc., can be made by combining these primitives.


SUSPEND AND RESUME

Sometimes workflows need to pause execution while waiting for something outside the agent's control — a human approval, a slow external API, or a long-running async process.


Because the third party can take arbitrarily long to respond, you don't want to keep a running process.

Instead, you want to persist the state of the workflow, and have some function that you can call to pick up where you left off.

Let's diagram out a simple example with Mastra's `.suspend()` and `.resume()` functions:



To handle suspended workflows, you can watch for status changes and resume execution when ready:



Here's an example of what creating a workflow with suspend and resume looks like in Mastra.

We start with **steps**, which are the building blocks of workflows. Create a step using `createStep`:


```

// Create a step with defined input/output schemas and execution logic
const inputSchema = z.object({
  inputValue: z.string(),
});

const myStep = createStep({
  id: "my-step",
  description: "Does something useful",
  inputSchema,
  outputSchema: z.object({
    outputValue: z.string(),
  }),
  // Optional: Define the resume schema for step resumption
  resumeSchema: z.object({
    resumeValue: z.string(),
  }),
  // Optional: Define the suspend schema for step suspension
  suspendSchema: z.object({
    suspendValue: z.string(),
  }),
  execute: async ({
    inputData,
    mastra,
    getStepResult,
    getInitData,
    runtimeContext,
  }) => {
    const otherStepOutput = getStepResult(step2);
    const initData = getInitData<typeof inputSchema>(); // typed as the input schema
    variable (zod schema)
    return {
      outputValue: `Processed: ${inputData.inputValue}, ${initData.startValue}
(runtimeContextValue: ${runtimeContext.get("runtimeContextValue")})`,
    };
  },
});

```

Then create a workflow using `createWorkflow`:

```

// Create a workflow with defined steps and execution flow
const myWorkflow = createWorkflow({
  id: "my-workflow",
  // Define the expected input structure (should match the first step's inputSchema)
  inputSchema: z.object({
    startValue: z.string(),
  }),
  // Define the expected output structure (should match the last step's outputSchema)
  outputSchema: z.object({
    result: z.string(),
  }),
  steps: [step1, step2, step3], // Declare steps used in this workflow
})
  .then(step1)
  .then(step2)
  .then(step3)
  .commit();

// Register workflow with Mastra instance
const mastra = new Mastra({
  vnext_workflows: {
    myWorkflow,
  },
});

// Create a run instance of the workflow
const run = mastra.vnext_getWorkflow("myWorkflow").createRun();

```

Once you've defined a workflow, run it like so:



```
// Create a run instance
const run = myWorkflow.createRun();

// Start the workflow with input data
const result = await run.start({
  inputData: {
    startValue: "Initial data",
  },
});

// Access the results
console.log(result.steps); // All step results
console.log(result.steps["step-id"].output); // Output from a specific step

if (result.status === "success") {
  console.log(result.result); // The final result of the workflow, result of the last
  // step (or '.map()' output, if used as last step)
} else if (result.status === "suspended") {
  const resumeResult = await run.resume({
    steps: result.suspended[0], // there is always at least one step id in the suspended
    // array, in this case we resume the first suspended execution path
    resumeData: {
      /* user input */
    },
  });
} else if (result.status === "failed") {
  console.error(result.error); // only exists if status is failed, this is an instance of
  // Error
}
```

And that's it. We've just built an AI agent workflow in Mastra that branches, pauses execution, and then resumes, all in a couple dozen lines of code.

In production, make sure your workflow state is persisted to a durable store — a suspended workflow that lives only in memory won't survive a server restart.


STREAMING UPDATES

One of the keys to making LLM applications feel fast and responsive is showing users what's happening, in incremental stages if necessary, while the model is working.

Back in 2024, I was trying to plan a trip to Hawaii for myself and my partner. I tried two early reasoning models side by side: OpenAI's o1 pro (left tab) and Deep Research (right tab).

The o1 pro just showed a spinning “reasoning” box for three minutes-no feedback, just waiting. Deep Research, on the other hand, immediately asked me for details (number of people, budget, dietary needs), then streamed back updates as it found restaurants and attractions.

The contrast was clear: the model that streamed better felt way snappier and kept me in the loop the whole time.



Left: 01 pro (less good). Right: Deep Research (more good)

Streaming step completion

The most common thing to stream is the model's output, showing tokens as they're generated. But you can also stream updates from each step in a multi-step workflow or agent pipeline, like when an agent is searching, planning, and summarizing in sequence.

Streaming within steps

Even when you're inside a function with a defined return type (eg a workflow step), you should be writing to a stream as new results come in. Here's how would with Mastra:

```
import { createStep } from '@mastra/core/workflows'
import { z } from 'zod'

export const testStep = createStep({
  execute: async ({ inputData, mastra, writer }) => {
    const { city } = inputData

    const testAgent = mastra?.getAgent('testAgent')
    const stream = await testAgent?.stream('What is the weather in London?')

    await stream!.textStream.pipeTo(writer!)

    return {
      value: await stream!.text,
    },
  },
})
```

Streaming tool calls

One thing that's improved: in March 2026, most major APIs now support streaming tool call results natively.

Speed matters

Like in web apps and mobile games, perceived response time is crucial in user engagement. There are no commercially successful agents that feel slow.

And on a personal note: after having to reschedule my vacation a couple times, my partner and I finally did go in June 2025.

While hiking in Volcanoes National Park, we turned a corner and saw Mount Kilauea shooting lava hundreds of feet in the air.

Nature sometimes delivers a different kind of streaming update....

PART V

**RETRIEVAL-AUGMENTED
GENERATION (RAG)**

RAG 101

RAG lets agents ingest user data and synthesize it with their global knowledge base to give users high quality responses.

Here's how it works.

Chunking: You start by taking a document (although we can use other kinds of sources as well) and chunking it. We want to split the document into bite-sized pieces for search.

Embedding: After chunking, you'll want to embed your data – transform it into a vector, or an array of 1536 values between 0 and 1, representing the meaning of the text.

We do this with LLMs, because they make the embeddings much more accurate; OpenAI has an API for this, there are other providers like Voyage or Cohere.

You need to use a vector DB which can store these vectors and do the math to search on them. You can use pgvector, which comes out of the box with Postgres.

Indexing: Once you pick a vector DB, you need to set up an index to store your document chunks, represented as vector embeddings.

Querying: Okay, after all that setup, you can now query the database!

Under the hood, you'll be running an algorithm that compares your query string to all the chunks in the database and returning the most similar ones. The most popular algorithm is called "cosine similarity".

The implementation is similar to a geospatial query searching latitude/longitude, except the search goes over 1536 dimensions instead of two.

You can use other algorithms as well.

Reranking: Optionally, after querying, you can use a reranker. Reranking is a more computationally expensive way of searching the dataset. You can run it over your results to improve the ordering (but it would take too long to run it over the entire database).

Synthesis: Finally, you pass your results as context into an LLM, along with any other context you want to include, and it can synthesize an answer to the user.

CHOOSING A VECTOR DATABASE

One of the biggest questions people having around RAG is how they should think of a vector DB.

There are multiple form factors of vector databases:

1. A feature on top of open-source databases (**pgvector** on top of Postgres)
2. Standalone open-source (**Chroma**)
3. Standalone hosted cloud service (**Turbopuffer**).
4. Hosted by an existing cloud provider (**Cloudflare Vectorize**).

Our take is that unless your use-case is exceptionally specialized, the vector DB feature set is mostly commoditized.

Back in 2023, VC funding drove a huge explosion in vector DB companies.

In 2026, many of these companies are doing layoffs (hint: any vector DB company that's done layoffs is also struggling with adoption).

Today, teams report wanting to prevent infra sprawl (yet *another* service to maintain). Our recommendation:

- If you're already using Postgres for your app backend, pgvector is a fine choice.
- If you're spinning up a new project, Chroma or Turbopuffer are great defaults.
- If your cloud provider has a managed vector DB service, start with that.

SETTING UP YOUR RAG PIPELINE

RAG combines the strengths of pre-trained LLMs with the ability to retrieve and incorporate specific and relevant information from your own curated datasets in order to generate more accurate and detailed (and sometimes more up-to-date) responses. Setting up a RAG pipeline involves loading data, creating embeddings, indexing, retrieving relevant information based on a query, and then generating a response using the LLM with the retrieved context.

Chunking

Chunking is the process of breaking down large documents into smaller, manageable pieces for processing.

The key thing you'll need to choose here is a **strategy** and an **overlap window**. Good chunking balances context preservation with retrieval granularity.

Strategies for doing the actual chunking itself include recursive, character-based, token-aware, and format-specific

corpus splitting like Markdown, HTML, JSON, or LaTeX. Mastra supports all of these.

Embedding

Embeddings are numerical representations of text that capture semantic meaning. These vectors allow us to perform similarity searches. Mastra supports multiple embedding providers like OpenAI and Cohere, with the ability to generate embeddings for both individual chunks and arrays of text.

Upsert

Upsert operations allow you to insert or update vectors and their associated metadata in your vector store. This operation is essential for maintaining your knowledge base, combining both the embedding vectors and any additional metadata that might be useful for retrieval.

Indexing

An index is a data structure that optimizes vector similarity search. When creating an index, you specify parameters like dimension size (matching your embedding model) and similarity metric (cosine, euclidean, dot product). This is a one-time setup step for each collection of vectors.

Querying

Querying involves converting user input into an embedding and finding similar vectors in your vector store. The basic query returns the most semantically similar chunks to your

input, typically with a similarity score. Under the hood, this is a bunch of matrix multiplication to find the closest point in n -dimensional space (think about a geo search with lat/long, except in 1536 dimensions instead).

The most common algorithm that does this is called *cosine similarity* (although you can use others instead).

Hybrid Queries with Metadata. Hybrid queries combine vector similarity search with traditional metadata filtering. This allows you to narrow down results based on both semantic similarity and structured metadata fields like dates, categories, or custom attributes.

Reranking

Reranking is a post-processing step that improves result relevance by applying more sophisticated scoring methods. It considers factors like semantic relevance, vector similarity, and position bias to reorder results for better accuracy.

It's a more computationally intense process, so you typically don't want to run it over your entire corpus for latency reasons — you'll typically just run it on a code example.

Code Example

Here's some code using Mastra to set up a RAG pipeline. Mastra includes a consistent interface for creating indexes, upserting embeddings, and querying, while offering their own unique features and optimizations, so while this example uses PgVector, you could easily use another DB instead. This illustrates a common pattern for building context-aware AI responses with proper fallback handling.

```

import { Mastra } from "@mastra/core";
import { MDocument, PgVector } from "@mastra/rag";
import { embedMany, embed } from "ai";
import { openai } from "@ai-sdk/openai";

// Initialize document and create chunks
const doc = MDocument.fromText("Your text content here...");
const chunks = await doc.chunk({
  strategy: "recursive",
  size: 512,
  overlap: 50,
});

// Generate embeddings
const { embeddings } = await embedMany({
  values: chunks.map(chunk => chunk.text),
  model: openai.embedding("text-embedding-3-small"),
});

// Initialize vector store and Mastra
const pgVector = new
PgVector({process.env.POSTGRES_CONNECTION_STRING});

// Store embeddings
await pgVector.createIndex("embeddings", 1536);
await pgVector.upsert(
  "embeddings",
  embeddings,
  chunks?.map(chunk => ({ text: chunk.text }))
);

// Query example
const query = "insert query here";
const { embedding } = await embed({
  value: query,
  model: openai.embedding("text-embedding-3-small"),
});

// Retrieve similar chunks
const results = await pgVector.query("embeddings", embedding);
const relevantContext = results
  .map(result => result?.metadata?.text)
  .join("\n\n");

// Generate response
const completion = await openai("gpt-4o-mini").generate(
  `Please answer the following question:
  ${query}

  Based on this context: ${relevantContext}
  If the context lacks sufficient information, please state that
  explicitly.`
);

console.log(completion.text);

```

This code:

1. Uses OpenAI's GPT-4o-mini model to generate a response
2. Constructs a prompt that includes a user query and relevant context
3. Instructs the model to explicitly state if there's insufficient information
4. Logs the generated text response

Note: There are advanced ways of doing RAG: using LLMs to generate metadata, using LLMs to refine search queries; using graph databases to model complex relationships. These may be useful for you, but start by setting up a working pipeline and tweaking the normal parameters — embedding models, rerankers, chunking algorithms — first.

ALTERNATIVES TO RAG

Great, now you know how RAG works. But does it matter? Or, put like a Twitter edgelord, is RAG dead?

Not yet, we think. But there are some simpler approaches you should reach for first.

Give your agent search tools

Instead of pre-parsing all your documents, start by giving your agent a relevant set of search tools to search through its corpus of data.

One of our investors built a variety of tools to query her website, and then bundled them into a MCP server she could give to an agent.

She then asked the agent about her favorite restaurants (it recommended Flour + Water in San Francisco) and her favorite portfolio companies (it said she likes all of her companies equally).

Let your agent run code

Opus 4.6 and GPT 5.4 are really good at writing code. Consider solving your search problem by giving your model a code writing tool and running it in a loop. (*What would Claude Code do?*)

Feed the model full context

Gemini has 1m tokens now. Maybe you should just load all the relevant content directly into the context.

The downsides: cost and “context rot” (recall can drop far before hitting context window limits.)

Extract entities and relationships

Use a model to extract rich semantic information (section references, entity names, etc), and then throw those into the context window as well.

Conclusion

We’re engineers. We can over-engineer things.

Unless you’re sure this is a pure document processing problem, build an agent first. Use RAG as a fallback.

PART VI

MULTI-AGENT SYSTEMS

MULTI-AGENT 101

Think about a multi-agent systems like a specialized team, like marketing or engineering. Different agents work together with their own specialized role to accomplish more complex tasks.

If you've used agentic coding tools Claude Code or Replit Agent, you've actually already been using a multi-agent system.


One agent works with you to plan / architect your code. After you've mapped it all out, you next work with a "code manager" agent that passes instructions to a code writer, executes the resulting code in a sandbox, and passes any errors back to the code writer.

Each of these agents has different memories, different system prompts, and access to different tools.

We often joke that designing a multi-agent system involves a lot of skills used in designing an organization. You group related tasks into a job description. You might give creative tasks to one person and analytical tasks to another.

You want to think about network dynamics. Is it better for three specialized agents to gossip among themselves? Or

feed their output back to a manager agent who can make a decision?



Multi-agent systems with subagents let you break down complex tasks into manageable pieces. And designs are fractal. A hierarchy is just a supervisor of supervisors.

But when you're building, start with the simplest version first.

Let's break down some of the patterns.

AGENT SUPERVISOR & SUBAGENTS

Agent supervisors are specialized agents that coordinate and manage other “subagents”.

The most straightforward way to do this is to pass in the other agents wrapped as tools.

For example, in a content creation system, a publisher agent might supervise both a copywriter and an editor:

```
const publisherAgent = new Agent({  
  name: "publisherAgent",  
  instructions: "You are a publisher agent that coordinates content creation. First call the copywriter for initial content, then the editor for refinement.",  
  model: {  
    provider: "ANTHROPIC",  
    name: "claude-3-5-sonnet-20241022",  
  },  
  tools: { copywriterTool, editorTool },  
});
```


CONTROL FLOW

When building complex AI applications, you need a structured way to manage how agents think and work through tasks. Just as a project manager wouldn't start coding without a plan, agents should establish an approach before diving into execution.

Just like how it's common practice for PMs to spec out features, get stakeholder approval, and only then commission engineering work, you shouldn't expect to work with agents without first aligning on what the desired work is.

We recommend engaging with your agents on architectural details first — and perhaps adding a few checkpoints for human feedback in their workflows.

WORKFLOWS AS TOOLS

Hopefully, by now, you're starting to see that all multi-agent architecture comes down to which primitives you're using and how you're arranging them.

It's particularly important to remember this framing when trying to build more complex tasks into agents.

Let's say you want your agent(s) to always run 3 separate tasks sequentially. You may want to turn each of those tasks into individual workflows, so you can stipulate a workflow's order of steps and provide more structure.

Each of these workflows can then be passed along as *tools* to the agent(s).

PARALLELIZED TOOL CALLS

In most of the workflows we've covered, tool calls are sequential by default.

Parallel tool calling lets the model batch those calls into a single step. If the results don't depend on each other, they can run simultaneously. You can imagine opening several URLs in parallel to complete research faster, or running several coding subagents in parallel to test different approaches.

Over the last several months, all three major providers have added support. OpenAI added it earliest; Anthropic has extended the idea the furthest.

Today, Claude models enable parallel tool calls by default. You can disable this if your workflow requires strict sequential execution — for instance, when earlier tool results need to gate what gets called next.

COMBINING THE PATTERNS

By this point, you'll notice there are different architectures. If your agentic coding tool has a plan mode and a build mode, the agents embody different steps in a workflow.

If your customer support agent is delegating to a billing workflow, a product docs subagent, or escalating to a human, you can see the value of subagents and subworkflows.

The lesson: each agent primitive can be arranged to suit your project's desired control flow.

A note on A2A

Some people still associate multi-agent with A2A. In reality, A2A is a niche protocol (Google Trends data shows 20 searches for MCP for every A2A search). Multi-agent has moved beyond A2A.

PART VII

OBSERVABILITY & EVALS

TRACING

There are two uniquely hard parts of building AI applications, that are typically the barrier for teams shipping agents into production: accuracy and token cost.

The answer for both: observability.

Let's dive in.

Why accuracy and token cost matter

Accuracy. AI applications are built on top of non-deterministic models. Teams that have shipped agents into production typically building systems to detect regressions and look at production data every day. Agents can regress while still returning 200 OK.

Token cost. Agents burn tokens. Opus 4.6 and GPT 5.4 can run in the background, in loops, generating multiple potential responses. But this can be costly.

We know several startups that raised \$5m seed rounds, went viral, and ended up with a \$500k token bill the next

month. You can make something people want, but if your tokens cost 10x your revenue, you still have a big problem.

Observability

Observability is a word that gets a lot of airplay, but since its meaning has been largely diluted and generalized by self-interested vendors let's go to the root.

The initial term was popularized by Honeycomb's Charity Majors in the late 2010s to describe the quality of being able to visualize application traces.

To debug a function, it would be *really nice* to be able to see the input and output of every function it called. And the input and output of every function *those* functions called. (And so on, and so on, turtles all the way down.)

This kind of telemetry is called a *trace*, which is made up of a tree of *spans*. (Think about a nested HTML document, or a flame chart.)

The standard format for traces is known as Open-Telemetry, or OTel for short.

When monitoring vendors began supporting tracing, each had a different spec — there was no portability. Lightstep's Ben Sigelman helped create the common Otel standard, and larger vendors like Datadog (under duress) began to support Otel.

Visualizing traces

There's a large number of observability vendors, both older backend and AI-specific ones, but the UI patterns converge:

you're developing. Here's a screenshot of how this looks in Mastra Studio:

EVALS

While traditional software tests have clear pass/fail conditions, AI outputs are non-deterministic — that is, an LLM can take the same input and produce different outputs. (You’ve seen how ChatGPT and Claude offer the option to regenerate the results of your query as soon as they finish spitting it out, right?)

Evals help bridge this gap by providing quantifiable metrics for measuring agent quality.

Think about evals sort of like including, say, performance testing in your CI pipeline. There’s going to be some randomness in each result, but on the whole and over time there should be a correlation between application performance and test results.

When writing evals, it’s important to think about what *exactly* you want to test.

There are different kinds of evals just like there are different kinds of tests.

Unit tests are easy to write and run but might not capture the behavior that matters; end-to-end tests might

capture the right behavior but they might also be more flaky.

Similarly, if you're building a RAG pipeline or a structured workflow, you may want to test each step along the way — and then after that test the behavior of the system as a whole.


Seeing evals over time

It's also nice to be able to see your evals (more on evals later) in a cloud environment with your team.

For each of their evals, people want to see a side-by-side comparison of how the agent responded versus what was expected.

They want to see the overall score on each PR (to ensure there aren't regressions), and the score over time, and to filter by tags, run date, and so on.

Eval UIs tends to look like this:



A sample evaluation screen

. . .

Offline vs. online evals

There are two modes teams run evals in.

Offline evals run against a fixed dataset before deploys. These tend to catch regressions.

Online evals run against live production traffic. You sample real requests, score them, and watch metrics over time. These can be noisy (your inputs vary over time). But real users break your agents in new and exciting ways.

Most teams start offline. Then teams add online evals once their agent is in production.

Building your eval dataset


There are three ways to get a dataset of test cases.

Hand-curate. This is a good place to start because it forces you to think clearly about what good looks like. You may want to work with a subject-matter expert if you're working in an unfamiliar domain.

Generate them synthetically. This is fast, but you should check the output; models tend to generate easy cases.

Mine production logs. This only works when you're *in* production, but it's the highest-signal source.

A mature eval dataset is a mix of all three. Like normal code, as you find bugs, fix them and then add test cases to prevent future regressions.



Datasets in Mastra Studio

LLM-as-judge for textual evals

For textual evals, the most common pattern is using a second model as the scorer.

You pass the output — and optionally the original input and retrieved context — to a judge model with a rubric. It returns a score and sometimes an explanation.

It works well. It scales. It handles cases where there's no single right answer. It also does well judging metrics like hallucination and faithfulness (whether responses accurately represent provided context).

It's not perfect: judges tend to prefer longer answers even when they're not better. Consider picking a judge from a different model family as your agent; a judge can favor outputs that sound like itself.

```
import { Agent } from '@mastra/core/agent'
import { createFaithfulnessScorer } from '@mastra/evals/scorers/prebuilt'

export const evaluatedAgent = new Agent({
  scorers: {
    relevancy: {
      scorer: createFaithfulnessScorer({ model: 'openai/gpt-5.1' }),
    },
  },
})
```

Adding a LLM-as-judge for faithfulness in Mastra

Classification or labeling evals

Classification or labeling evals help determine how accurately a model tags or categorizes data based on predefined categories (e.g., sentiment, topics, spam vs. not spam).

This can include broad labeling tasks (like recognizing document intent) or fine-grained tasks (like identifying specific entities aka entity extraction).

Tool calling evals

Tool usage evals for agents measure how effectively a model or agent calls external tools or APIs to solve problems.

For example, just like you would write `expect(Fn).toBeCalled` in the JavaScript testing framework Jest, you would want to see similar functions for agent tool use.

Multi-turn evals

Most textual evals are built around single-turn interactions. But agents are multi-turn by nature.

A multi-turn eval runs the agent through a full conversation or task sequence and grades the whole thing. Did it maintain context across turns? Did it recover when a tool failed? Did it stay on task after a tangent?

Task completion

For agents doing things in the world, the most important eval is the simplest one: did it finish the job?

Did the support agent resolve the ticket? Did the coding

agent pass the tests? Did the research agent return a usable answer?

Prompt engineering evals

Prompt engineering evals explore how different instructions, formats, or phrasings of user queries impact an agent's (non-deterministic) results performance.

They look at both the **sensitivity** of the agent to prompt variations (whether small changes produce large differences in results) and the **robustness** to adversarial or ambiguous inputs.

All things “prompt injection” fall in this category.

A/B testing

Leaders of larger consumer and developer tools AI companies like Perplexity and Replit joke that they rely more on A/B testing of user metrics than evals per se. This works if you have enough traffic that degradation in agent quality is quickly visible!

Human data review

In addition to automated tests, high-performing AI teams regularly review production data. Typically, the easiest way to do this is to view traces which capture the input and output of each step in the pipeline.

Many correctness aspects (e.g., subtle domain knowledge, or an unusual user request) can't be fully captured by rigid assertions, but human eyes catch these nuances.

Observability & evals are even more important than you think

This might be relatively long chapters, but observability is one of the most important pieces when shipping agents into production.

Teams can spend months staring closely at their observability tools as they move from prototype to production.

We've written a follow-up book covering much more on these topics: **Patterns For Building AI Agents**. We encourage you to read it, especially as you're pushing your agent into production.

PART VIII

**DEVELOPMENT &
DEPLOYMENT**

LOCAL DEVELOPMENT

Agent development typically falls into two different categories: building the frontend vs. backend.

Building an agentic web frontend

Web-based agent frontends tend to share a few characteristics: they're built around a chat interface, stream to a backend, autoscroll, and display tool calls.

We discussed the importance of streaming in an earlier chapter. Agentic interfaces tend to use a variety of different transport options like request/response, server-sent events, webhooks and web sockets, all to feed the sense of real-time interactivity.

There are a few frameworks we see speeding up development here, especially during the prototype phase: Assistant UI, Copilot Kit, and Vercel's AI SDK UI.

(And many agents are based on other platforms like WhatsApp, Slack, mobile apps, email, or voice and don't have a web frontend.)

It's important to note that while agentic frontends can be

powerful, the full agent logic generally can't live client-side in the browser for security reasons — it would leak your API keys to LLM providers.

Building an agent backend


So it's the backend where we typically see most of the complexity.

When developing AI applications, it's important to see what your agents are doing, make sure your tools work, and be able to quickly iterate on your prompts.

Some things that we've seen be helpful for local agent development:

- **Agent chat interface:** Test conversations with your agents in the browser, seeing how they respond to different inputs and use their configured tools.
- **Workflow visualizer:** Seeing step-by-step workflow execution and being able to suspend/resume/replay
- **Agent/workflow endpoints:** Being able to curl agents and workflows on localhost (this also enables using eg Postman)
- **Tool playground:** Testing any tools and being able to verify inputs / outputs without needing to invoke them through an agent.
- **Tracing & evals:** See inputs and outputs of each step of agent and workflow execution, as well as eval metrics as you iterate on code.

Mastra's local dev environment, which we call Mastra Studio, lives at `localhost:4111`:

The image is a screenshot of the 'Mastra Playground' interface, specifically the 'Agents' tab. On the left, there is a sidebar with 'Agents' selected, and below it, 'Tools' and 'Workflows' are visible. The main area displays a table of agents. The table has five columns: 'Name', 'Instruction', 'Provider', 'Model', and 'Action'. There are five rows of agents listed: 'Specify Agent', 'Weather Agent', 'cat-one', 'Code Review Agent', and 'Chef Agent'. Each row includes a small icon to the left of the name and a 'Chat with agent' button to the right of the model name. At the bottom left of the interface, there is a breadcrumb trail: 'Mastra AI > Docs > OPENAI'.

Name	Instruction	Provider	Model	Action
 Specify Agent	You are a helpful assistant that can make recipes based on products.	OPENAI.CHAT	gpt-4o	 Chat with agent
 Weather Agent	You are a helpful weather assistant that provides accurate weather information. Your primary function is to help users get weather ...	OPENAI.CHAT	gpt-4o	 Chat with agent
 cat-one	You are a feline expert with comprehensive knowledge of all cat species, from domestic breeds to wild big cats. As a feline cat s...	OPENAI.CHAT	gpt-4o	 Chat with agent
 Code Review Agent	An agent that reviews code repositories and provides feedback.	OPENAI.CHAT	gpt-4o	 Chat with agent
 Chef Agent	You are MICHEL, a practical and experienced home chef who helps people cook great meals with whatever ingredients they have ex...	OPENAI.CHAT	gpt-4o	 Chat with agent

In Mastra Studio, you can interact with your agents in real-time, observe their decision-making process, review prompts and responses, and debug function calls so you can iterate quickly on prompts, workflows, and integration logic.

DEPLOYMENT

Most teams are wrapping their agents in a server, wrapping the server in a Docker container, and scaling that.

The core agent loop

The core agent loop is long-running and stateful. Sure, you can model each “turn” as a separate request/response cycle.

But you still have to think about streaming a response back to the user, rather than simply sending a single JSON blob.

Agentic workflows

Relatively speaking, similar to the workloads on durable execution engines like Temporal and Inngest. But they are still tied to a specific user request.

Agents can run on serverless platforms, but teams need to work through dealing with timeouts from long-running processes, bundle size issues, etc.

Using a managed platform

The agent teams sleeping the soundest at night are the ones we see who figure out how to run their agents using auto-scaling managed services.

Teams using container services like AWS EC2, Digital Ocean, or equivalent seem to be all right as long as they have a B2B use case that won't have sudden usage spikes.

PART IX

CODING AGENTS

AGENTS BUILDING AGENTS

In early 2026, as the agentic coding experience has gotten better, developers have become increasingly addicted to Claude Code.

Venture capitalist Nikunj Kothari joked about “token anxiety” — people leaving Silicon Valley parties early so they could get back to running their agents.

Earlier, in 2025, coding IDEs like Cursor allowed users to install MCP servers to provide agents with more tools and knowledge. (We ourselves shipped an open-source MCP server with access to our docs).


On web-based platforms

Meanwhile, vibe coding platforms like Replit, Lovable, and Emergent have crossed \$100M and \$200M in revenue and shipped products letting their users build agents.

One example: in September 2025, Replit’s Agent 3 shipped support for vibe coding agents with natural language. Unlike drag-and-drop builders, the agents gener-

ated by Replit's agent are fully code-based, and can be exported and run on other platforms.

Replit chose to use Mastra for the framework of the generated agents, and exposed a simple UI for the agents that helps users visualize the work that's written.



Using coding agents to build agents opens the door for non-technical users that could only previously build using no-code tools with inherent platform lock-in. It also lets developers build agents even more quickly.

SANDBOXES & FILESYSTEMS

Any agent that writes and executes code needs somewhere safe to run it.

You probably don't want to run agent-generated code directly on your servers — this creates both resource contention and security risks.

Sandboxes provide an isolated environment where code runs separately from your application via containers, micro-VMs, or intercepted system calls.

In 2026, there's a category of dedicated platforms like E2B, Daytona, and Modal. Cloud providers like Cloudflare and Vercel also offer sandbox primitives.

Ephemeral vs. stateful sandboxes

For one-off tasks, you want an ephemeral sandbox — it can spin up, runs some code, and disappear.

For longer-horizon tasks, you may want a stateful sandbox that persists between calls — your agent can install dependencies, write files, and return to the same environment later.

Filesystems

Most sandbox platforms expose a filesystem API so your agent can read and write files just like it would on a local machine. Whether that filesystem survives between sessions depends on the platform — ephemeral sandboxes wipe it on teardown, stateful ones persist it.

Some platforms also let you snapshot the filesystem and restore to it later. This is particularly useful for evals: run your agent against a known starting state, check what changed, roll back, and run again.

Managing latency

Sandboxes add latency. A 150ms startup is fine for a coding agent that runs for minutes, but noticeable if you're spinning up a new sandbox on every tool call.

Some teams strike a balance: create one sandbox per agent session, reuse it across all tool calls, then tear it down when the session ends.

PART X

EVERYTHING ELSE

MULTIMODAL

One way to think about multimodality (images, video, voice) in AI is to map their dates of origin on various platforms.

Consider the Internet: it supported text from its origin in the 1970s, but images and video weren't supported until the web browser (1992), and voice not until 1995.

Voice and video didn't become popular until Skype (2003) and YouTube (2005), with greater bandwidth and connection speeds.

Or think about social networks: all the early ones, like MySpace (2002), Facebook (2004), and Twitter (2008), were primarily text-based.

Image-based social media didn't become popular until Instagram (2010) and Snapchat (2013), and video-based social media until TikTok (2017).


In AI, then, it's little wonder that multi-modal use-cases are a bit younger and less mature. Like on earlier platforms, they're trickier to get right, and more computationally complex.

Image Generation

March 2025 brought the invention of Ghibli-core —think soft colors, dreamy backgrounds, and those iconic wide-eyed characters.

People had been playing with Midjourney, Stable Diffusion, and others for a while. Transposing photos into specific styles with a single prompt, though, was a significant leap forward in consumer-grade image-generation.

People uploaded selfies or old photos, added a prompt, and instantly got back an anime version that looked straight out of “Spirited Away.”



The Mastra cofounders (Shane, Abhi and Sam) at a basketball game

This wasn't just a niche thing; the Ghibli trend took over social feeds everywhere. The official (Trump) White House account joined the fray by (controversially) tweeting out a Ghibli-fied picture of a detained immigrant.

More broadly, the Ghibli moment showed vitality for the digital art use case — image gen for what was something

that is a confluence between storyboard, a character sketch, and environment concepts.

Use Cases

In terms of people using image gen for products, there are a few use-cases.

In marketing and e-commerce, product mockups on varied backgrounds and rapid ad creative generation without photoshoots and in various form factors. “Try-on” image models allow people to swap out the human model but keep the featured clothing item.

The third use-case for image gen has been in video game and film production. Image gen has allowed for asset prototyping, including portraits, textures, and props, as well as scene layout planning via rough “sketch to render” flows.

Put in web development terms, this gives the fidelity of a full design with the effort/skill of a wireframe.

Last, there are more NSFW use-cases. These don’t tend to be venture-fundable, but at least according to the Silicon Valley gossip mills, quite a few of the more risqué use-cases print money — if you can find a payment processor that will take your business.

Voice

The key modalities in voice are speech-to-text (STT), text-to-speech (TTS), and speech-to-speech, also known as realtime voice.

Most voice agents built in 2024 and early-to-mid 2025 used a speech-to-text (STT) and text-to-speech (TTS) pipeline. They use one model to translate input voice to text,

another model to generate response text, and then translate the response text into an audio response.

Mastra offers some voice capabilities for agents. Here's an example of listening; you could follow this up with `agent.speak()` to reply.

```
import { Agent } from "@mastra/core/agent";
import { OpenAIVoice } from "@mastra/voice-openai";

const agent = new Agent({
  name: 'Agent',
  instructions: `You are a helpful assistant with
voice capabilities.`,
  model: openai('gpt-4o'),
  voice: new OpenAIVoice();
});

const audioStream =
fs.createReadStream('/path/to.mp3')

const text = await agent.listen(audioStream)

// Hey! How are ya!
```

But capabilities are evolving quickly; end-to-end speech-to-speech models like OpenAI's `gpt-realtime` are increasingly deployed in production voice agents.

In order to do that, you could train a model that specifically takes voice tokens as input, and responds with voice tokens as output. That's known as “realtime voice”, but it's proved challenging.

For one thing, it's difficult to train such models; the information density of audio is only 1/1000 of text, so these models take significantly more input data to train and cost more to serve.

Second, these models still struggle with turn-taking, known in the industry as “voice activity detection”. When talking, humans interrupt each other constantly using both visual and emotional cues.

But voice models don’t have these cues, and have to deal with both computational and network latency. When they interrupt too early, they cut people off; when they interrupt too late, they sound robotic.

While these products make great demos, there are not too many companies using realtime voice in production.

Video

For most of 2025, we were waiting for video to have their Studio Ghibli moment where they can accurately represent characters in input and replay them in alternate settings.

Now in early 2026, tools like Sora 2, Veo 3.1 (Google), Runway Gen-4.5, and Kling 3.0 are in active production use. Veo 3.1 and Seedance 2.0 have achieved character consistency.

We’re starting to see them show up in social media ads, corporate training, film concept visualization, YouTube automation, and multilingual content.

WHAT'S NEXT

This book was intentionally written as a getting-started guide. We have a sequel, *Patterns for Building AI Agents*, that dives more into productionizing them.

Patterns includes more on context engineering, creating high-quality eval suites by working with domain experts and analyzing failure modes, plus more on agent security.

The agent space is moving incredibly quickly.

We don't have a crystal ball, but from our vantage point as a prominent agent framework, here's what we see:

- **Reasoning models will continue to get better.**

Claude Opus 4.6 introduced "adaptive thinking," where the model dynamically allocates reasoning effort based on problem complexity. GPT-5.4 offers explicit reasoning effort controls (low through xhigh) so you can tune latency/depth per request. As the model does more planning, our job may shift from choreographer to architect.

- **We'll make progress on agent learning.** Agents emit traces, but right now the feedback loop to improve their performance runs through their human programmers.
- **Synthetic evals are becoming standard practice.** Generating evals automatically from production traces, with human approval, has moved from experimental to mainstream. People aren't going to hand code evals in the era of Claude Code.
- **Agent & MCP security will continue to be a challenge.** As agents ingest more data and take more actions, prompt injection attacks will multiply. And MCP's rapid adoption has created a large and fast-growing attack surface.
- **The eternal September of AI will continue.** Every month brings new developers who haven't learned how to write good prompts or how to do context engineering. Meanwhile, the rapid pace of model updates means even established teams are constantly adapting their implementations.

In a field where the ground shifts constantly, we're all perpetual beginners. To build something enduring, you have to stay humble.

